

ECE 253 Homework 3

Name: Sheng-Wei Chang

PID: A53317226

Academic Integrity Policy: Integrity of scholarship is essential for an academic community. The University expects that both faculty and students will honor this principle and in so doing protect the validity of University intellectual work. For students, this means that all academic work will be done by the individual to whom it is assigned, without unauthorized aid of any kind. By including this in my report, I agree to abide by the Academic Integrity Policy mentioned above.

```
In [1]: import numpy as np
import cv2
import matplotlib.pyplot as plt
from scipy.signal import convolve
import warnings
warnings.filterwarnings('ignore')
```

Problem 1. Canny Edge Detection

```
In [2]: def CannyEdgeDetection(img, t):
n, m = img.shape
# Smoothing
kg = np.array([[2, 4, 5, 4, 2],
               [4, 9, 12, 9, 4],
               [5, 12, 15, 12, 5],
               [4, 9, 12, 9, 4],
               [2, 4, 5, 4, 2]]) / 159
img = convolve(img, kg, mode='same')

# Finding Gradients
# x: left to right, y: top to bottom
kx = np.array([[-1, 0, 1],
               [-2, 0, 2],
               [-1, 0, 1]])
ky = np.array([[-1, -2, -1],
               [0, 0, 0],
               [1, 2, 1]])
Gx = convolve(img, kx, mode='same')
Gy = convolve(img, ky, mode='same')
Gm = np.sqrt(Gx**2 + Gy**2)
Gd = np.arctan2(Gy, Gx)

# Non-maximum Suppression
Gd = np.round(Gd/(np.pi/4))
nms = np.zeros_like(Gm)
for i in range(n):
    for j in range(m):
        if Gd[i,j] == 3 or Gd[i,j] == -1:
            if Gm[i,j] >= Gm[min(n-1,i+1),max(0,j-1)] and Gm[i,j] >= Gm[max(0,i-1),min(n-1,j+1)]:
                nms[i,j] = Gm[i,j]
        elif Gd[i,j] == 2 or Gd[i,j] == -2:
            if Gm[i,j] >= Gm[min(n-1,i+1),j] and Gm[i,j] >= Gm[max(0,i-1),j]:
                nms[i,j] = Gm[i,j]
        elif Gd[i,j] == 1 or Gd[i,j] == -3:
            if Gm[i,j] >= Gm[min(n-1,i+1),min(m-1,j+1)] and Gm[i,j] >= Gm[max(0,i-1),max(0,j-1)]:
                nms[i,j] = Gm[i,j]
        else:
            if Gm[i,j] >= Gm[i,min(m-1,j+1)] and Gm[i,j] >= Gm[i,max(0,j-1)]:
                nms[i,j] = Gm[i,j]

# Thresholding
output = np.where(nms>t, 1, 0)

return Gm, nms, output, Gd
```

In [3]:

```
# Read image
img = cv2.imread('geisel.jpg')

# Convert the image into grayscale before doing histogram equalization
geisel = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# Histogram equalization
t = 50
Gm, nms, output, Gd = CannyEdgeDetection(geisel, t)

# plot
plt.figure(figsize=(18,12))
plt.subplot(2,2,1)
plt.title('Original Image')
plt.imshow(geisel, cmap='gray')
plt.axis('off')
plt.subplot(2,2,2)
plt.title('Original gradient magnitude image')
plt.imshow(Gm, cmap='gray')
plt.axis('off')
plt.subplot(2,2,3)
plt.title('Image after NMS')
plt.imshow(nms, cmap='gray')
plt.axis('off')
plt.subplot(2,2,4)
plt.title('The final edge image after thresholding.')
plt.imshow(output, cmap='gray')
plt.axis('off')
plt.show()

print('The value of te I used to produce the final edge image is %d.' %t)
```

Original Image



Original gradient magnitude image






Image after NMS



The final edge image after thresholding.



The value of te I used to produce the final edge image is 50.

Problem 2. Butterworth Notch Reject Filtering in Frequency Domain

In [4]:

```
def padding(img):
    n, m = img.shape
    return np.pad(img, (((512-n)//2, 512-n-(512-n)//2), ((512-m)//2, 512-m-(512-m)//2)), 'constant')

def fft(img):
    out = np.fft.fftshift(np.fft.fft2(img))
    return out, np.log(np.abs(out))

def ifft(img):
    f_ishift = np.fft.ifftshift(img)
    return np.abs(np.fft.ifft2(f_ishift))

def butterworth_notch_reject_filter(K, n, D0, uk, vk):
    x_axis = np.linspace(-256, 255, 512)
    y_axis = np.linspace(-256, 255, 512)
    [u,v] = np.meshgrid(x_axis, y_axis)

    H = np.ones((512, 512))
    for k in range(K):
        Dp = np.sqrt((u-uk[k])**2 + (v-vk[k])**2)
        Dn = np.sqrt((u+uk[k])**2 + (v+vk[k])**2)
        H *= (1/(1+(D0/Dp)**(2*n))) * (1/(1+(D0/Dn)**(2*n)))
    return H

def plot_results(img1, img2, img3, img4, img5):
    plt.figure(figsize=(18, 12))
    plt.subplot(2,3,1)
    plt.imshow(img1, cmap='gray')
    plt.title('Original image')
    plt.axis('off')
    plt.colorbar()
    plt.subplot(2,3,2)
    plt.imshow(img2, cmap='gray')
    plt.title('FT log-magnitude')
    plt.axis('off')
    plt.colorbar()
    plt.subplot(2,3,4)
    plt.imshow(img3, cmap='gray')
    plt.title('Butterworth Filter in frequency domain')
    plt.colorbar()
    plt.subplot(2,3,5)
    plt.imshow(img4, cmap='gray')
    plt.title('FT log-magnitude times a Butterworth filter')
    plt.axis('off')
    plt.colorbar()
    plt.subplot(2,3,6)
    plt.imshow(img5, cmap='gray')
    plt.title('Filtered image')
    plt.axis('off')
    plt.colorbar()

    plt.show()
```

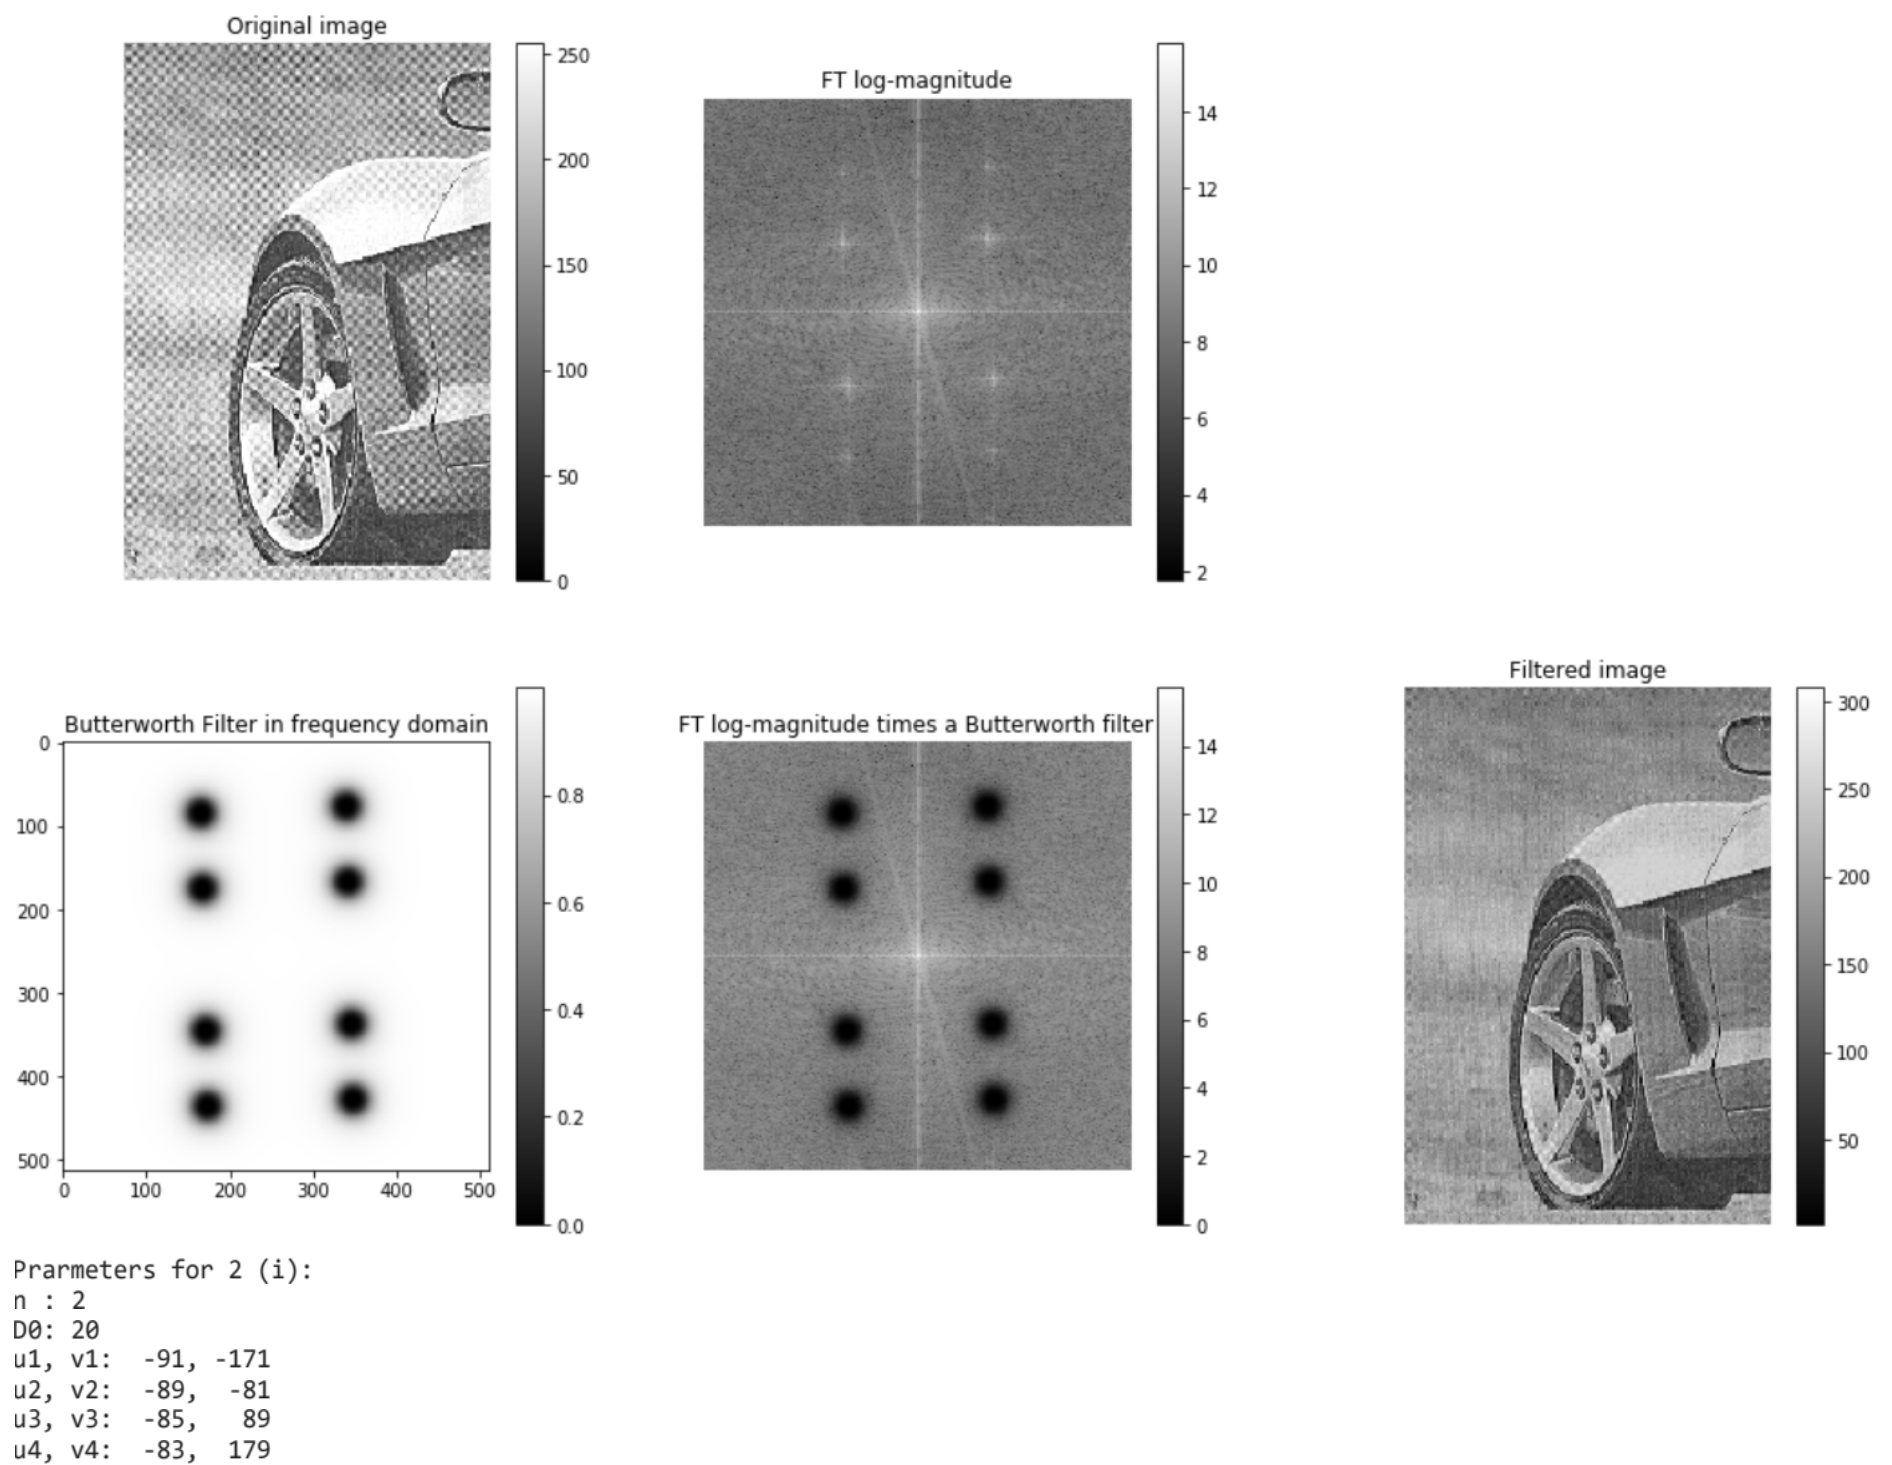
In [5]:

```
# Car
img = cv2.imread('Car.tif')
car = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
car_ft, car_ft_log = fft(padding(car))

K = 4
n = 2
D0 = 20
uk = np.array([165, 167, 171, 173]) - 256
vk = np.array([ 85, 175, 345, 435]) - 256

H = butterworth_notch_reject_filter(K, n, D0, uk, vk)
car_ft_n = car_ft_log * H
car_filtered = ifft(car_ft * H)
plot_results(car, car_ft_log, H, car_ft_n, car_filtered[133:379,172:340])

# 10 parameters for 2(i): n, D0, u1, v1, ..., u4, v4
print('Parameters for 2 (i):')
print('n : %d' % n)
print('D0: %d' % D0)
print('u1, v1: %4d, %4d' % (uk[0], vk[0]))
print('u2, v2: %4d, %4d' % (uk[1], vk[1]))
print('u3, v3: %4d, %4d' % (uk[2], vk[2]))
print('u4, v4: %4d, %4d' % (uk[3], vk[3]))
```



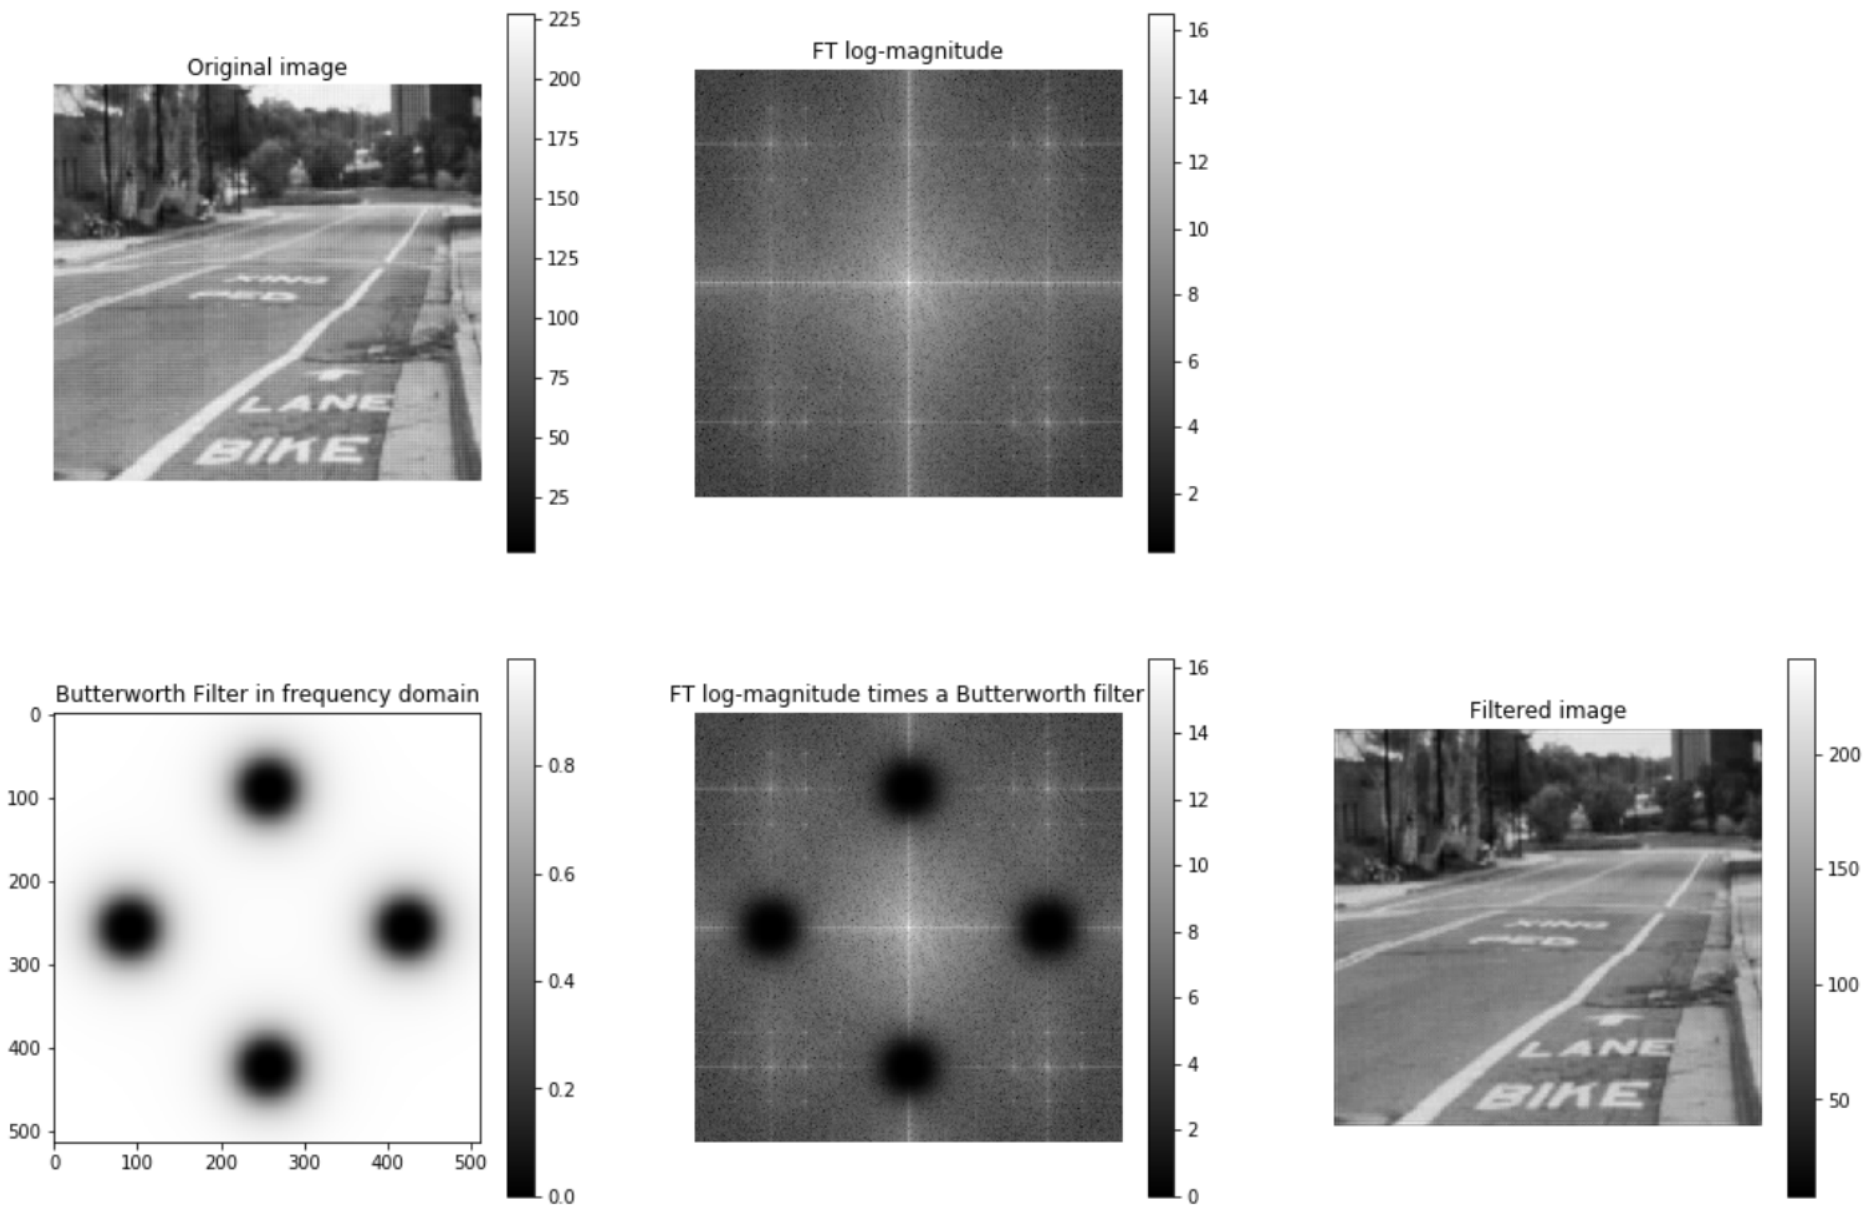
In [6]:

```
# Street
img = cv2.imread('Street.png')
street = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
street_ft, street_ft_log = fft(padding(street))

K = 2
n = 2
D0 = 40
uk = np.array([ 90, 256]) - 256
vk = np.array([256, 90]) - 256

H = butterworth_notch_reject_filter(K, n, D0, uk, vk)
street_ft_n = street_ft_log * H
street_filtered = ifft(street_ft * H)
plot_results(street, street_ft_log, H, street_ft_n, street_filtered[90:422,76:435])

# 6 parameters for 2(ii): n, D0, u1, v1, u2, v2
print('Prarmeters for 2 (ii):')
print('n : %d' % n)
print('D0: %d' % D0)
print('u1, v1: %4d, %4d' % (uk[0], vk[0]))
print('u2, v2: %4d, %4d' % (uk[1], vk[1]))
```



Prarmeters for 2 (ii):
n : 2
D0: 40
u1, v1: -166, 0
u2, v2: 0, -166

Problem 3. PyTorch tutorial and questions

```
In [1]: import torch
import torchvision
import torchvision.transforms as transforms
```

```
In [2]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda:0

```
In [3]: transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified
Files already downloaded and verified

```
In [4]: import matplotlib.pyplot as plt
import numpy as np

# functions to show an image
def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



bird car car ship

```
In [5]: import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
net = net.to(device)
```

```
In [6]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
In [7]: training_loss = []
for epoch in range(2): # Loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        # inputs, labels = data
        inputs, labels = data[0].to(device), data[1].to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
    if i % 2000 == 1999: # print every 2000 mini-batches
        print('[%d, %5d] loss: %.3f' % (epoch + 1, i + 1, running_loss / 2000))
        training_loss.append(running_loss / 2000)
        running_loss = 0.0

print('Finished Training')
```

```
[1, 2000] loss: 2.205
[1, 4000] loss: 1.877
[1, 6000] loss: 1.675
[1, 8000] loss: 1.587
[1, 10000] loss: 1.513
[1, 12000] loss: 1.443
[2, 2000] loss: 1.379
[2, 4000] loss: 1.353
[2, 6000] loss: 1.325
[2, 8000] loss: 1.299
[2, 10000] loss: 1.263
[2, 12000] loss: 1.242
Finished Training
```

```
In [8]: dataiter = iter(testloader)
images, labels = dataiter.next()

# print images
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



```
GroundTruth:   cat  ship  ship plane
```

```
In [9]: images, labels = data[0].to(device), data[1].to(device)
outputs = net(images)
```

```
In [10]: _, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]] for j in range(4)))
```

```
Predicted:  truck   cat   car   dog
```

```
In [11]: correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        # images, labels = data
        images, labels = data[0].to(device), data[1].to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
```

```
Accuracy of the network on the 10000 test images: 56 %
```

```
In [12]: class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        # images, labels = data
        images, labels = data[0].to(device), data[1].to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (classes[i], 100 * class_correct[i] / class_total[i]))
```

```
Accuracy of plane : 69 %
Accuracy of  car : 82 %
Accuracy of  bird : 21 %
Accuracy of  cat : 32 %
Accuracy of  deer : 61 %
Accuracy of  dog : 58 %
Accuracy of  frog : 65 %
Accuracy of horse : 59 %
Accuracy of  ship : 55 %
Accuracy of truck : 56 %
```

How many images and batches are used to train the network? ¶

```
In [13]: trainset
```

```
Out[13]: Dataset CIFAR10
  Number of datapoints: 50000
  Root location: ./data
  Split: Train
  StandardTransform
Transform: Compose(
  ToTensor()
  Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
)
```

There are 50000 images used to train the network. The batch size is 4, so there are 50000/4=12500 batches are used to train the network.

Do we normalize the images? What do we do in the example?

Yes, we normalize the images. The outputs of torchvision datasets are PILImage images of range [0, 1]. However, we need to normalize them to Tensors of normalized range [-1, 1], such that the mean is at 0. Thus, we use the formula below.


$$output = \frac{input - mean}{std}$$

We set both mean and std as 0.5. When the input is 0, the output is -1. When the input is 1, the output is 1. Also, this is a linear transformation.

The losses are dropping! Can you plot out the training loss?

```
In [14]: import matplotlib.pyplot as plt

plt.plot(training_loss)
plt.title('Training loss')
plt.grid(True)
plt.show()
```




Now the network is done training. Can you check some successful cases and some failure cases (show some images classified by the network)?

```
In [16]: dataiter = iter(testloader)
data = dataiter.next()
images, labels = data[0].to(device), data[1].to(device)

outputs = net(images)
_, predicted = torch.max(outputs.data, 1)

# print images
imshow(torchvision.utils.make_grid(images.cpu()))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
print('Predicted: ', ' '.join('%5s' % classes[predicted[j]] for j in range(4)))
print('The 1st and 4th images are successful cases. Their predicted labels are correct.')
print('The 2nd and 3rd image is a failure case.')
```



GroundTruth: cat ship ship plane
Predicted: cat car car plane
The 1st and 4th images are successful cases. Their predicted labels are correct).
The 2nd and 3rd image is a failure case.

Can you visualize the output of the 1st layer of CNN using one image from the training set?

```
In [22]: class FirstNN(nn.Module):
def __init__(self):
super(FirstNN, self).__init__()
self.conv1 = nn.Conv2d(3, 6, 5) # (number of input channels, number of output channels, kernel size)

def forward(self, x):
return self.conv1(x)

net_1st_layer = FirstNN().to(device)
```

```
In [23]: data = next(iter(testloader))
image, labels = data[0].to(device), data[1].to(device)
out = net_1st_layer(image)
print(out.shape) # (batch size, number of output channels, (32-5)+1=28)

torch.Size([4, 6, 28, 28])
```

```
In [28]: #image after convolution
image_1st_layer = out[0,:,:,:]

plt.figure(figsize = (12,5))
plt.subplot(2, 4, 1)
plt.imshow(image[0].cpu().numpy().transpose(1,2,0)*0.5+0.5)
plt.title('Input')
plt.axis('off')
for i in range(6):
plt.subplot(2, 4, i+2)
plt.imshow(image_1st_layer[i].detach().cpu().numpy())
plt.title('Output: Channel %d' % (i+1))
plt.axis('off')
plt.show()
```

